


# Progress Report

Kuei-Yueh (Clint) Ko

# Testing sample



```
N_POINTS = 10000  
DIM_GRID = 128  
mean = ([0.0, 0.0], [4.0, 4.0])  
prop = (0.5, 0.5)
```

# Benchmark of Scipy Rbf function

## Scipy Rbf('gaussian')

```
%%timeit
rbfi = Rbf(dat_pts[:, 0], dat_pts[:, 1], value, function='gaussian')
z_rbf_gauss = rbfi(x_c, y_c)
```

```
/local_data/env-py3/lib/python3.6/site-packages/scipy/linalg/basic.py:40: RuntimeWarning: scipy.linalg.solve
Ill-conditioned matrix detected. Result is not guaranteed to be accurate.
Reciprocal condition number/precision: 4.850221788728607e-20 / 1.1102230246251565e-16
RuntimeWarning)
```

15.7 s ± 26.7 ms per loop (mean ± std. dev. of 7 runs, 1 loop each)

## Scipy Rbf('multiquadric')

```
%%timeit
rbfi = Rbf(dat_pts[:, 0], dat_pts[:, 1], value, function='multiquadric')
z_rbf_multiq = rbfi(x_c, y_c)
```

```
/local_data/env-py3/lib/python3.6/site-packages/scipy/linalg/basic.py:40: RuntimeWarning: scipy.linalg.solve
Ill-conditioned matrix detected. Result is not guaranteed to be accurate.
Reciprocal condition number/precision: 9.993949597828553e-17 / 1.1102230246251565e-16
RuntimeWarning)
```

11.3 s ± 35.8 ms per loop (mean ± std. dev. of 7 runs, 1 loop each)

# Benchmark of my interpolation function

GPU (CPU input/output arrays); (threadsperblock = 32)

```
%%timeit  
get_weights[blockspersgrid, threadsperblock](grids, points, weights)
```

370 ms  $\pm$  17.9 ms per loop (mean  $\pm$  std. dev. of 7 runs, 1 loop each)

```
%%timeit  
get_weights[blockspersgrid, threadsperblock](grids, points, weights)  
z_test = np.matmul(weights, value.reshape(N_POINTS, -1))  
z_test = z_test.reshape(DIM_GRIDS, DIM_GRIDS)
```

844 ms  $\pm$  4.26 ms per loop (mean  $\pm$  std. dev. of 7 runs, 1 loop each)

**np.matmul ~ 844 - 370 = 474 (ms)**

- Threads per block = (32, 32)

# Benchmark of my interpolation function

**GPU (GPU input/output arrays); (threadsperblock = 32)**

```
%%timeit
get_weights[blockspergrid, threadsperblock](grids_device, points_device, weights_device)
weights = weights_device.copy_to_host()
```

452 ms  $\pm$  7.32 ms per loop (mean  $\pm$  std. dev. of 7 runs, 1 loop each)

```
%%timeit
get_weights[blockspergrid, threadsperblock](grids_device, points_device, weights_device)
weights = weights_device.copy_to_host()
z_test = np.matmul(weights, value.reshape(N_POINTS, -1))
z_test = z_test.reshape(DIM_GRIDS, DIM_GRIDS)
```

932 ms  $\pm$  489  $\mu$ s per loop (mean  $\pm$  std. dev. of 7 runs, 1 loop each)

**np.matmul ~ 932 - 452 = 480 (ms)**

**In this example, we do not gain a lot when the process of data copy from CPU to GPU is removed. I guess the reason is because this data copy process does not consume much time when comparing to calculating the weights**

- Threads per block = (32, 32)

# Note 01

## test skcuda and pycuda

```
import pycuda.autoinit
import pycuda.gpuarray as gpuarray
import numpy as np
import skcuda.linalg as linalg
import skcuda.misc as misc
linalg.init()
a = np.asarray(np.random.rand(4, 2), np.float32)
b = np.asarray(np.random.rand(2, 2), np.float32)
a_gpu = gpuarray.to_gpu(a)
b_gpu = gpuarray.to_gpu(b)
c_gpu = linalg.dot(a_gpu, b_gpu)
np.allclose(np.dot(a, b), c_gpu.get())
```

```
-----
OSError                                Traceback (most recent call last)
<ipython-input-263-bcef7324e87e> in <module>()
      4 import skcuda.linalg as linalg
      5 import skcuda.misc as misc
----> 6 linalg.init()
      7 a = np.asarray(np.random.rand(4, 2), np.float32)
      8 b = np.asarray(np.random.rand(2, 2), np.float32)

/local_data/env-py3/lib/python3.6/site-packages/skcuda/misc.py in init(allocator)
    175     global _global_cusolver_handle
    176     if not _global_cusolver_handle:
--> 177         from . import cusolver
    178         _global_cusolver_handle = cusolver.cusolverDnCreate()
    179

/local_data/env-py3/lib/python3.6/site-packages/skcuda/cusolver.py in <module>()
    49     break
    50 if _libcusolver == None:
--> 51     raise OSError('cusolver library not found')
    52
    53 class CUSOLVER_ERROR(Exception):
```

OSError: cusolver library not found

# Note 02

Threads per block = (64, 64) 

```
/usr/lib/python3/dist-packages/numba/cuda/cudadrv/driver.py in safe_cuda_api_call(*args)
    286         _logger.debug('call driver api: %s', libfn.__name__)
    287         retcode = libfn(*args)
--> 288         self._check_error(fname, retcode)
    289     return safe_cuda_api_call
    290

/usr/lib/python3/dist-packages/numba/cuda/cudadrv/driver.py in _check_error(self, fname, retcode)
    321         _logger.critical(msg, _getpid(), self.pid)
    322         raise CudaDriverError("CUDA initialized before forking")
--> 323         raise CudaAPIError(retcode, msg)
    324
    325     def get_device(self, devnum=0):
```

**CudaAPIError:** [1] Call to cuLaunchKernel results in CUDA\_ERROR\_INVALID\_VALUE